

Simplified Implementation of DeePC

MAE 271D

Sanjit Sarda

June 7, 2025

1 Abstract

Data-Enabled Predictive Control (DeePC) is a data-driven optimal control method that enables predictive control without requiring an explicit model of the system. DeePC exploits the behavioral system framework to learn the system dynamics without explicitly characterizing the system through system identification. In addition to discussing implementation notes and challenges, this project also shows how to speed up the compute of DeePC without loss of generality.

2 Introduction

Model Predictive Control involves solving a finite-horizon optimization problem based on a predictive model of the system. This requires an accurate parametric model, which we can find through system identification, which can be difficult for many nonlinear systems, especially black boxes.

Data-Enabled Predictive Control (DeePC) solves this problem by offering a model-free framework. DeePC operates directly on measured input/output data, by leveraging the behavioral systems framework. The system is represented within an implicit model stored within Hankel matrices constructed from historical data. This representation effectively captures the system's behavior including non-linearities by capturing rolling window information from past and present data, allowing modeling without explicit system identification.

3 Assumptions

This implementation of DeePC in this report is based on a 2D double integrator system. The assumptions underlying the simulation and control setup are outlined below:

3.1 System Dynamics

The system is modeled as a continuous-time double integrator with the following state-space representation:

$$\dot{x}(t) = Ax(t) + Bu(t), \quad y(t) = Cx(t)$$

where the state vector $x \in \mathbb{R}^4$ consists of position and velocity in two dimensions, and the control input $u \in \mathbb{R}^2$ represents acceleration commands. For a double integrator:

$$A = \begin{bmatrix} \mathbf{0} & I_n \\ \mathbf{0} & \mathbf{0} \end{bmatrix}, B = \begin{bmatrix} \mathbf{0} \\ I_n \end{bmatrix}, C = \begin{bmatrix} I_n & \mathbf{0} \end{bmatrix}, D = \mathbf{0}$$

The system is discretized using zero-order hold with a sampling period $dt = 1.0$, resulting in discrete-time matrices A_d, B_d .

3.2 Constraints

Box constraints are imposed on both the system state and inputs:

- Position and velocity components of the state are constrained to $[-10, 10]$.
- Control inputs (accelerations) are constrained to $[-2, 2]$.

3.3 Disturbance Model

Additive Gaussian noise is injected to the output of the system at each time step to simulate measurement disturbances. The covariance matrix was selected to be

$$\sigma = \begin{bmatrix} 0.1 & 0.05 \\ 0.05 & 0.04 \end{bmatrix}$$

And two means, μ were selected as $\mathbf{0}$ and $\mathbf{0.1}$ to test its performance on steady state error.

3.4 Data Collection

Offline input/output data is collected by applying random inputs spanning the range of allowable inputs. For a Linear System, this was found to be sufficient information, since a carefully designed persistently exciting input sequence.

Online Data Collection is collected by using an MPC to drive the system from a pre determined state to the desired initial state. This is designed to emulate online data collection in a real system, where the system will likely be manually controlled before allowing DeePC to autopilot and plan a trajectory.

3.5 Simplifications

To focus on DeePC's core mechanics, the following simplifications are made:

- The full system state is observable and used directly for output. This includes position and velocity data
- The system structure (double integrator) is known but not exploited—DeePC operates entirely from data.

- The dynamics are linear and time-invariant.
- Offline data is collected disturbance free.

4 Problem Formulation

DeePC very simply reformulates the optimal control problem by replacing the dynamics constraints with the implicit model constraints.

4.1 Traditional MPC Formulation

In MPC, given a discrete-time linear system:

$$x_{k+1} = Ax_k + Bu_k, \quad y_k = Cx_k$$

the objective is to minimize a cost over a finite horizon N , subject to dynamics and constraints:

$$\begin{aligned} \min_{u,x} \quad & \sum_{k=0}^{N-1} (\|x_k - x_{\text{ref}}\|_Q^2 + \|u_k\|_R^2) \\ \text{s.t.} \quad & x_{k+1} = Ax_k + Bu_k, \quad u_k \in \mathcal{U}, \quad x_k \in \mathcal{X} \end{aligned}$$

For this formulation, we are using the following Cost Parameters:

$$Q = \mathbf{I}, Q_f = 5\mathbf{I}, R = 2\mathbf{I}$$

A Receding Horizon Control is used, with 8 simulation steps, and a max horizon length of 4. Once fewer than 4 simulation steps are left, the horizon length becomes the remaining number of simulation steps.

4.2 DeePC Formulation

DeePC bypasses the need for the matrices A, B, C by using Hankel matrices derived from previously collected input/output data. These matrices implicitly encode the system behavior:

$$H_L(u) = \begin{bmatrix} u_1 & u_2 & \cdots & u_{T-L+1} \\ \vdots & \vdots & \ddots & \vdots \\ u_L & u_{L+1} & \cdots & u_T \end{bmatrix}$$

The Hankel matrix is a function of offline data and L . L represents the maximum context length available, and T represents how much offline data is within any given u . Therefore collecting more offline data results in a larger Hankel Matrix, and increasing the desired context length also increases the size of the Hankel Matrix.

We can then split the Hankel matrices into past and future data, where the past data is correlated with the *initial condition* context available to the controller, and the future data is correlated with the future of the controller.

Given:

- T_{ini} : length of initial trajectory,
- N : prediction horizon,
- $U_{\text{off}}, Y_{\text{off}}$: collected input/output data,

we construct the Hankel matrices:

$$\begin{bmatrix} U_p \\ U_f \end{bmatrix} = H_{T_{\text{ini}}+N}(U_{\text{off}}),$$

$$\begin{bmatrix} Y_p \\ Y_f \end{bmatrix} = H_{T_{\text{ini}}+N}(Y_{\text{off}})$$

where U_p, Y_p capture the past data used to infer the system's current state, and U_f, Y_f are used for prediction.

The DeePC optimization problem is formulated as:

$$\begin{aligned} \min_{g, u, y, \sigma_y} \quad & \sum_{k=0}^{N-1} \|y_k - r_k\|_Q^2 + \|u_k\|_R^2 + \lambda_g \|g\|_1 + \lambda_y \|\sigma_y\|_1 \\ \text{s.t.} \quad & \begin{bmatrix} U_p \\ Y_p \\ U_f \\ Y_f \end{bmatrix} g = \begin{bmatrix} u_{\text{ini}} \\ y_{\text{ini}} + \sigma_y \\ u \\ y \end{bmatrix} \\ & u_k \in \mathcal{U}, \quad y_k \in \mathcal{Y} \end{aligned}$$

Here, the implicit model is encoded within g .

For the sake of consistency and comparability, the same Cost parameters and Control Horizon is used here.

4.3 Regularization

The terms $\lambda_g \|g\|_1$ and $\lambda_y \|\sigma_y\|_1$ serve to regularize the optimization:

- λ_g penalizes large weights in the implicit model.
- λ_y accounts for measurement noise in the initial output trajectory.

5 Method

This implementation was written in Python using NumPy and CVXPY.

5.1 Simulation Framework

The controlled system is a two-dimensional double integrator with four states: positions and velocities in the x and y axes. The system is modeled in continuous time and discretized using SciPy's `cont2discrete` with a sampling period of 1.0s.

The simulation includes:

- Gaussian noise injection
- Hard box constraints on states and inputs.
- Step-by-step simulation logic in a `DoubleIntegrator` class.

5.2 Offline Data Collection Loop

A specified amount of data, $T = 200$ is collected. As previously mentioned, since this was a Linear System it was found any random set of inputs was sufficient to reliably estimate the system behavior.

5.3 Online Control Loop

A rolling horizon DeePC controller was implemented using the following steps:

1. After each time step, the most recent input/output data (length T_{ini}) is saved.
2. The DeePC optimization problem is solved using CVXPY.
3. The first control input from the solution is applied to the system.

5.4 Testing

Testing was performed by using an MPC to steer the system to a setpoint for online data collection and then using DeePC to steer the system to $\begin{bmatrix} 0 & 0 \end{bmatrix}$

The following starting points were used:

1. $\begin{bmatrix} 0 & 0 & 0 & 0 \end{bmatrix}$
2. $\begin{bmatrix} 0 & 0 & 0 & 5 \end{bmatrix}$
3. $\begin{bmatrix} 7 & -6 & -5 & -0.5 \end{bmatrix}$
4. $\begin{bmatrix} -3 & 2 & 0.5 & 0 \end{bmatrix}$
5. $\begin{bmatrix} 0 & 0 & -2.5 & -4 \end{bmatrix}$

5.5 Interpretability

There are two interesting ways to interpret DeePC.

5.5.1 Visualizing the Hankel Matrix

If you look at the Hankel matrix, you can see the diagonal structure in the matrix that shows it as a sliding window of data.




Figure 1: Hankel Matrix of the Offline Data Collected

5.5.2 System Identification

To demonstrate System Identification using a Hankel Matrix, and why DeePC works, a useful demonstration is to extract the discretized A_d and B_d matrices from the offline data.

Since the Hankel Matrix is a sliding Window of data, we can say

$$X_f = AX_p + BU_p$$

$$x_{k+1} = Ax_k + Bu_k$$

$$\text{So, } \begin{bmatrix} A & B \end{bmatrix} = X_f \begin{bmatrix} X_p \\ U_p \end{bmatrix}^\dagger$$

To run this with our offline data, and correctly collect the AB matrices, we need our output to be the whole state variable and the order to be 1, such that only past context is needed (no second derivatives in double integrator with pos and velocity state). Running this on the offline collected data, we get

$$\text{Estimated A matrix: } \begin{bmatrix} 1 & -0 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ -0 & 0 & 1 & 0 \\ -0 & 0 & -0 & 1 \end{bmatrix} \quad \text{Estimated B matrix: } \begin{bmatrix} 0.5 & -0 \\ 0 & 0.5 \\ 1 & 0 \\ -0 & 1 \end{bmatrix}$$

Which is exactly what we expect.

6 Results

First, we should look at the performance of an MPC controller on the system so that we have something to compare DeePC against.

These are the Hyperparameters to tune.

- **Order Estimate:** This controls the "context" window for the implicit model to understand the current state of the system.
- **T:** Length of Offline samples collected.

MPC Trajectories. Average Cost: 117.86

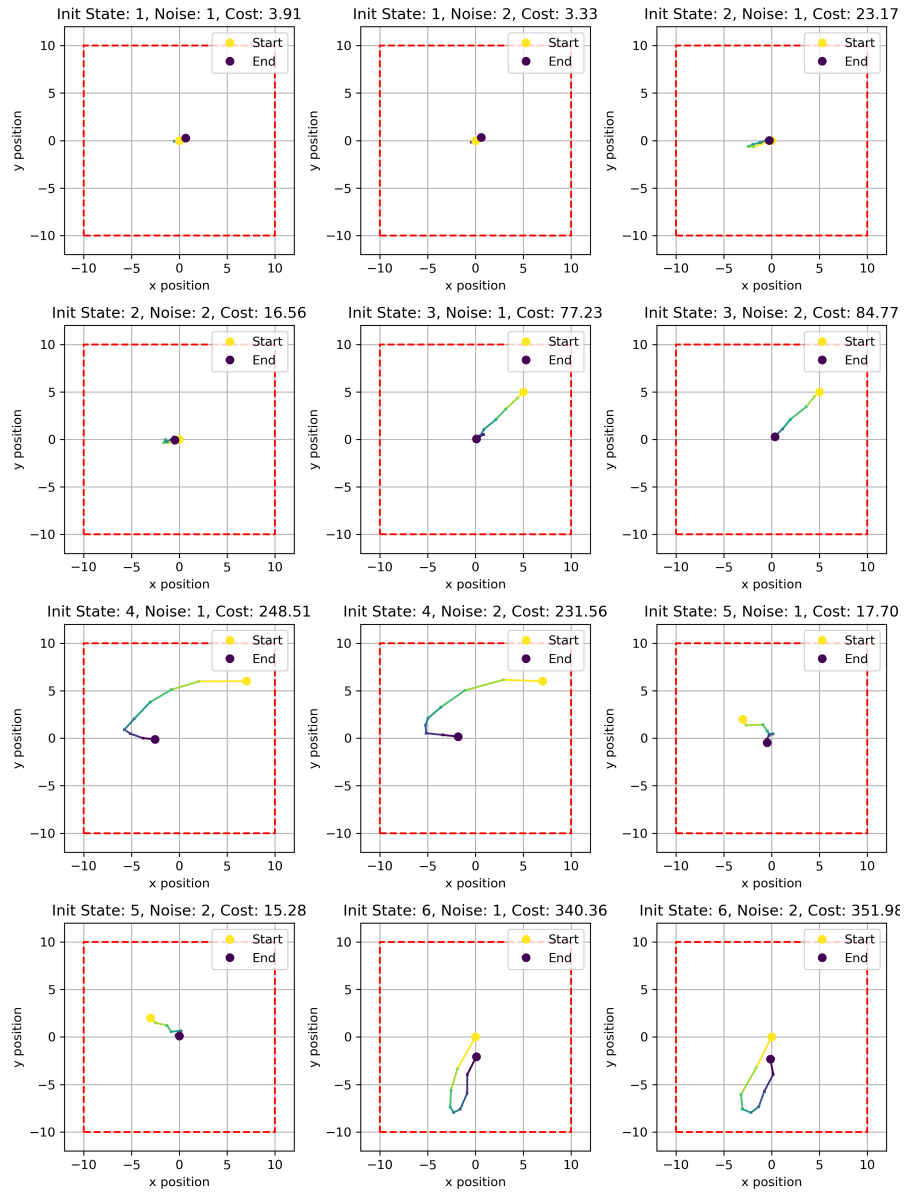


Figure 2: MPC performs fairly well, and optimizes over the the entire state.

- λ_g : Regulates the model size
- λ_y : Regulates the weight on the slack variable.

6.1 Hyperparameter Testing

Here are figures from Hyperparameter Testing with analysis

Other tests not mentioned: These tests dailed to converge:

- Heavily Over Estimating the Order: Using order = 10 completely failed.
- Under Regulating the Give: Without regulating the give, the value of sigma grows extremely high leading the model to be meaningless.

The following trends were observed

- **Order Estimate:** It is a good rule of thumb to match the correct estimate or go slightly above. If it fails then decrease the estimate order
- **T:** For a linear system more or less data does not affect performance too much
- λ_g : This did not affect the performance for the linear system, but this would be important for noisy data collection/ a non linear system
- λ_y : This is really important to keep high so that the model is actually constrained

6.2 Constraint Satisfaction

DeePC respected all input and state box constraints after the initial alignment phase, since the constraints are embedded within the optimization problem. Some constraint violations occured when the implicit model g did not well represent the model for some hyperparameters and during the offline data collection period.

6.3 Extension: Speeding up DeePC

The offline data collection leaves a lot of redundancy in the Hankel matrix, which leads to slow solve times. This can be fixed by using a Low Rank Approximation of the Hankel Matricies. This is done with SVD, and taking the full rank of the Hankel Matrix rather than an overdetermined system.

On average this computes trajectories over 50% faster for large datasets($T > 500$), without any noticable performance decrease.

Using a Low Rank Approximation speeds up the DeePC solver over 50% in the above tests. However there are two important notes to make:

Performance highly depends on the "rank" of the system. It is significantly easier to speed up a linear system than a system with a complex behavior, since the linear system has a low rank.

DeePC Trajectories, Over Estimate Lg=2.5, Ly=10000.0 Average Cost: 74.12

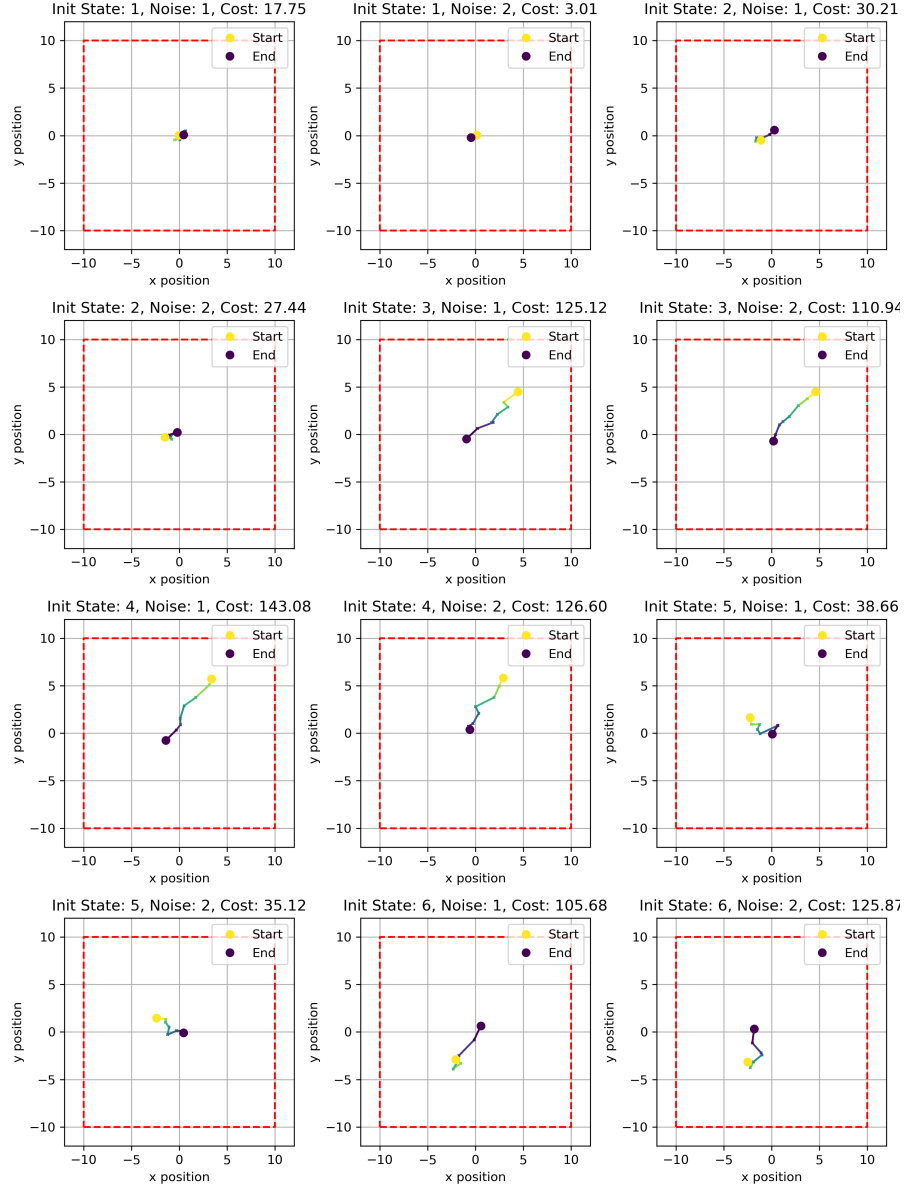


Figure 3: $\lambda_g = 2.5, \lambda_y = 1e4, \text{order} = 2, T = 100$ Performs fairly well

DeePC Trajectories, $T=100$, Est Order=1 $L_g=2.5$, $L_y=10000.0$ Avg Cost: 10115.6, Solve Time=0.02271s

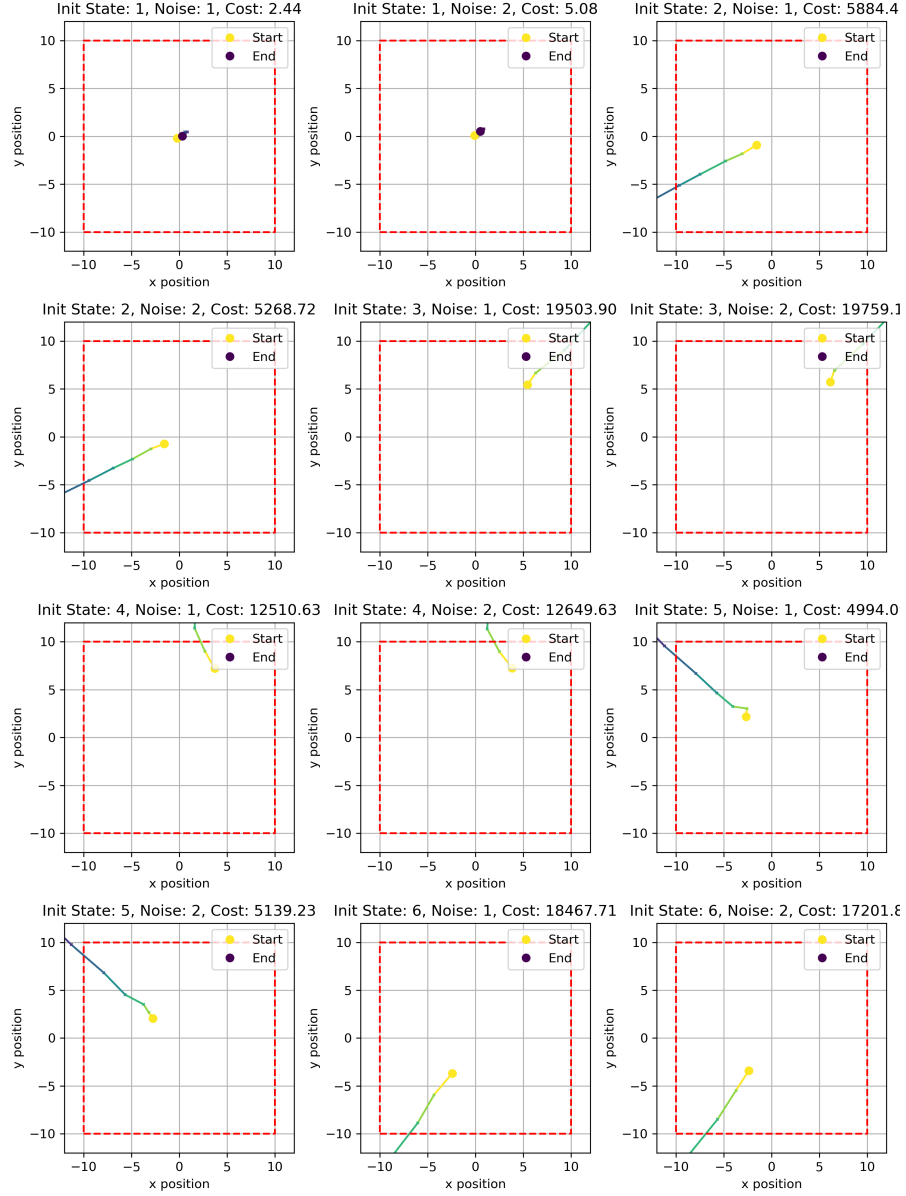


Figure 4: $\lambda_g = 2.5, \lambda_y = 1e4, \text{order} = 1, T = 100$ Completely collapses, since there is not enough data

DeePC Trajectories, $T=100$, Est Order=3 $L_g=2.5$, $L_y=10000.0$ Avg Cost: 72.3, Solve Time=0.02304s

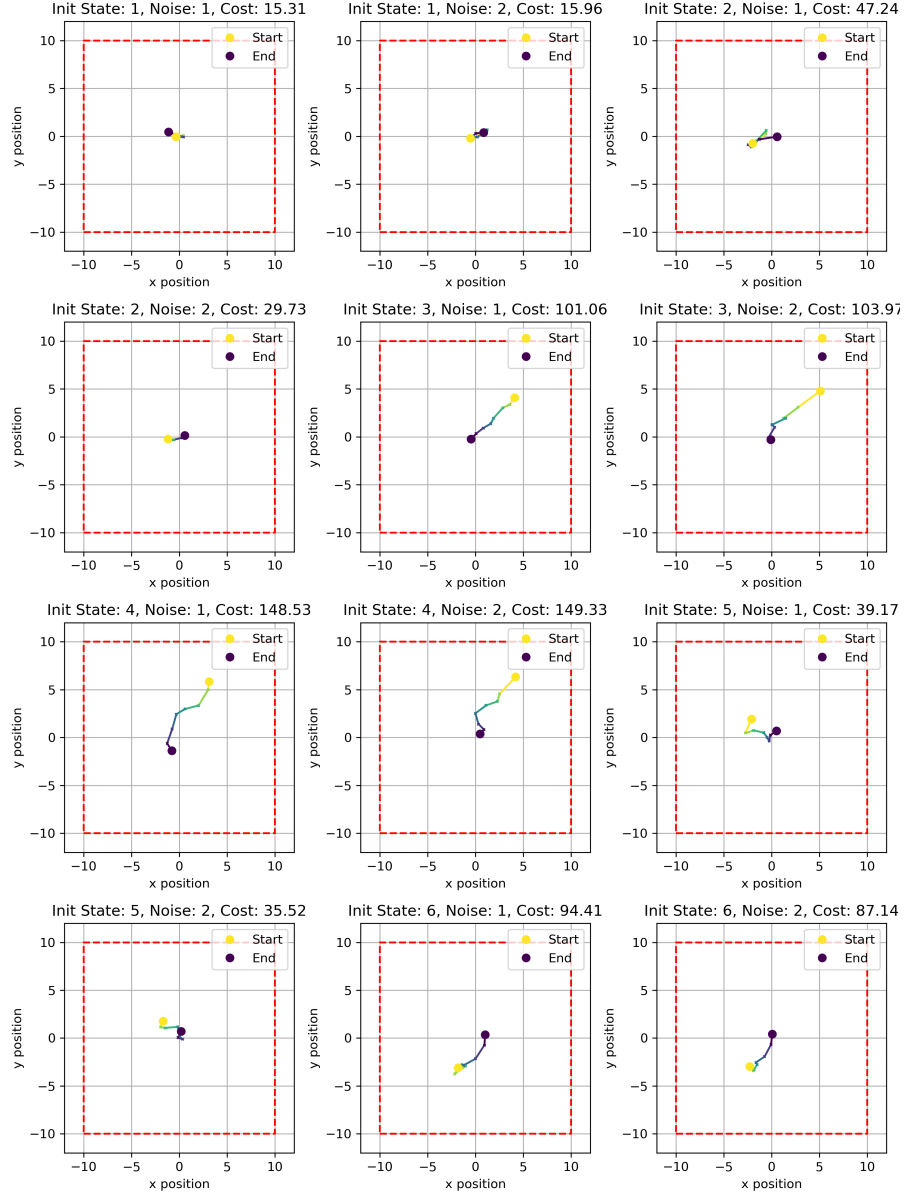


Figure 5: $\lambda_g = 2.5, \lambda_y = 1e4, \text{order} = 3, T = 100$ Slightly over estimating the order does not cause issues

DeePC Trajectories, T=25, Est Order=2 Lg=2.5, Ly=10000.0 Avg Cost: 79.8, Solve Time=0.01806s

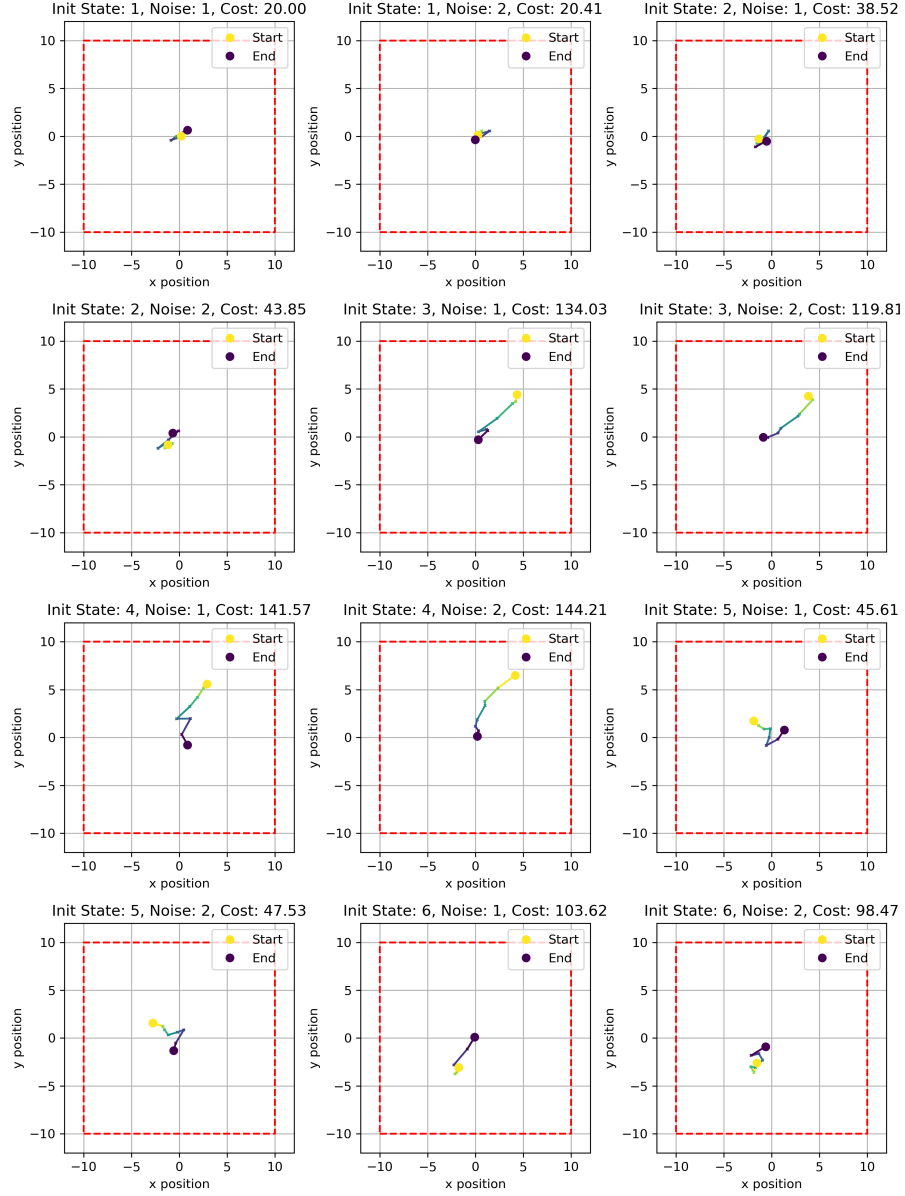


Figure 6: $\lambda_g = 2.5, \lambda_y = 1e4, \text{order} = 2, T = 25$ Performs fairly well, since it is a linear system, we do not actually need too much data

DeePC Trajectories, $T=100$, Est Order=2 $L_g=0.1$, $L_y=10000.0$ Avg Cost: 85.2, Solve Time=0.02922s

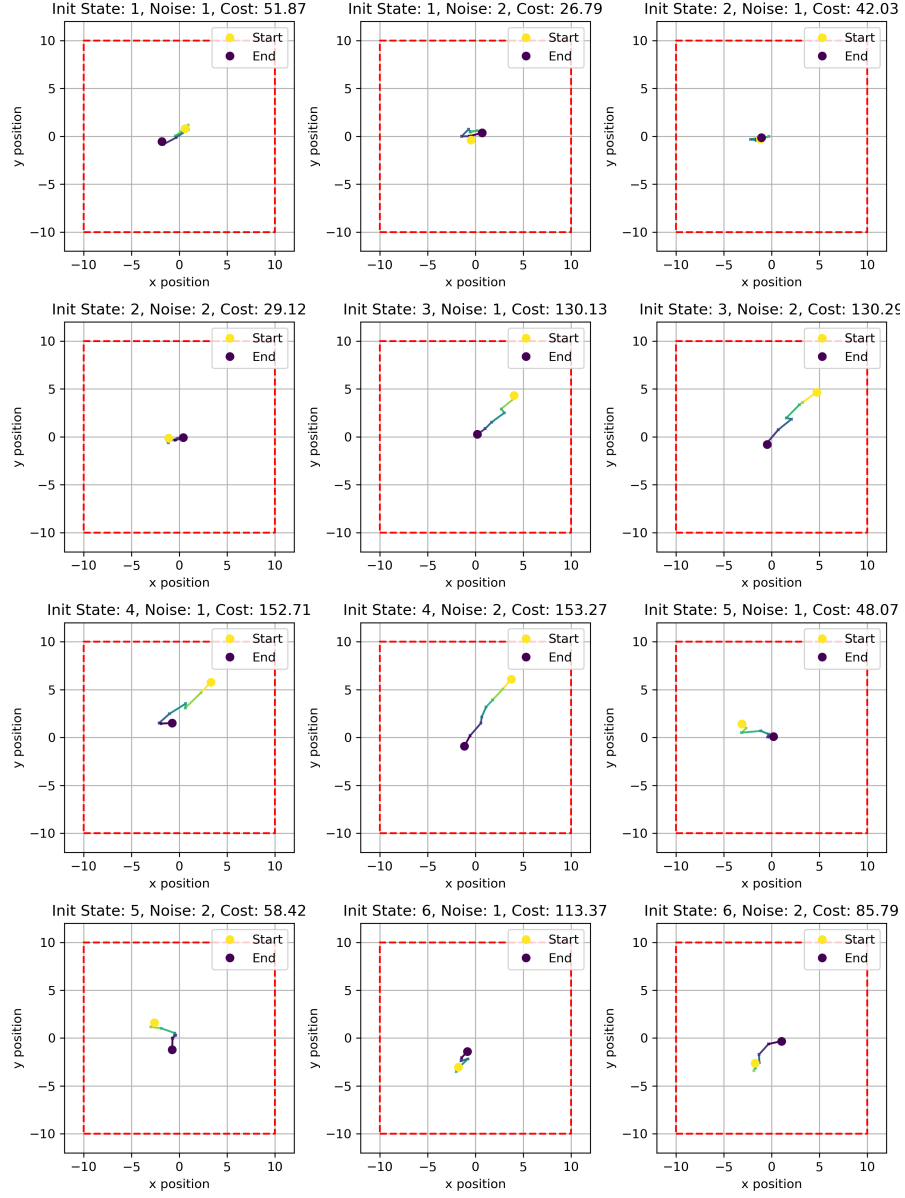


Figure 7: $\lambda_g = 0.1, \lambda_y = 1e4, \text{order} = 2, T = 100$ Since it is a linear system you do not need to regulate the model size too much.

DeePC Trajectories, T=100, Est Order=2 Lg=10, Ly=10000.0 Avg Cost: 75.3, Solve Time=0.03251s

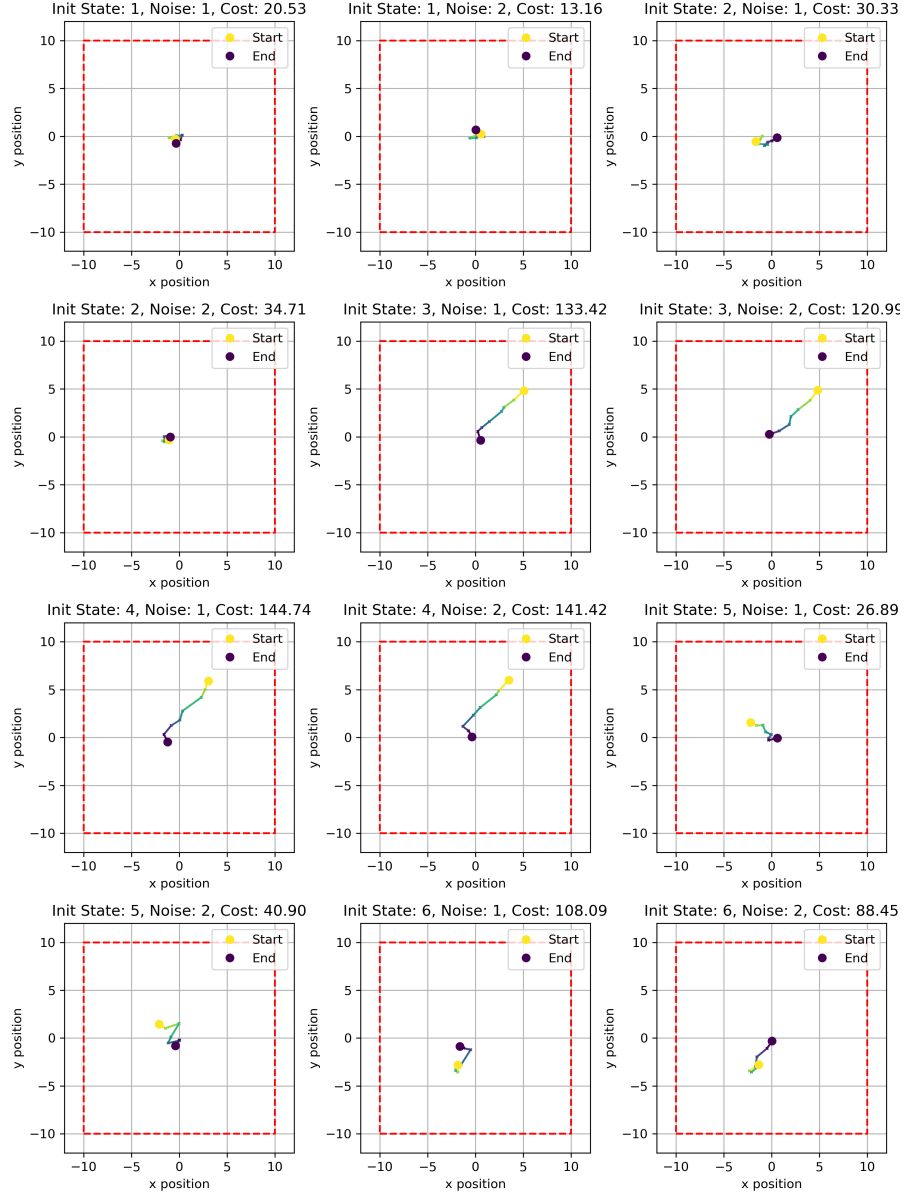


Figure 8: $\lambda_g = 10, \lambda_y = 1e4, \text{order} = 2, T = 25$ Heavily regulating the model size leads to adverse performance.

DeePC Trajectories, $T=100$, Est Order=2 $L_g=2.5$, $L_y=100000000.0$ Avg Cost: 1272.6, Solve Time=0.02620s

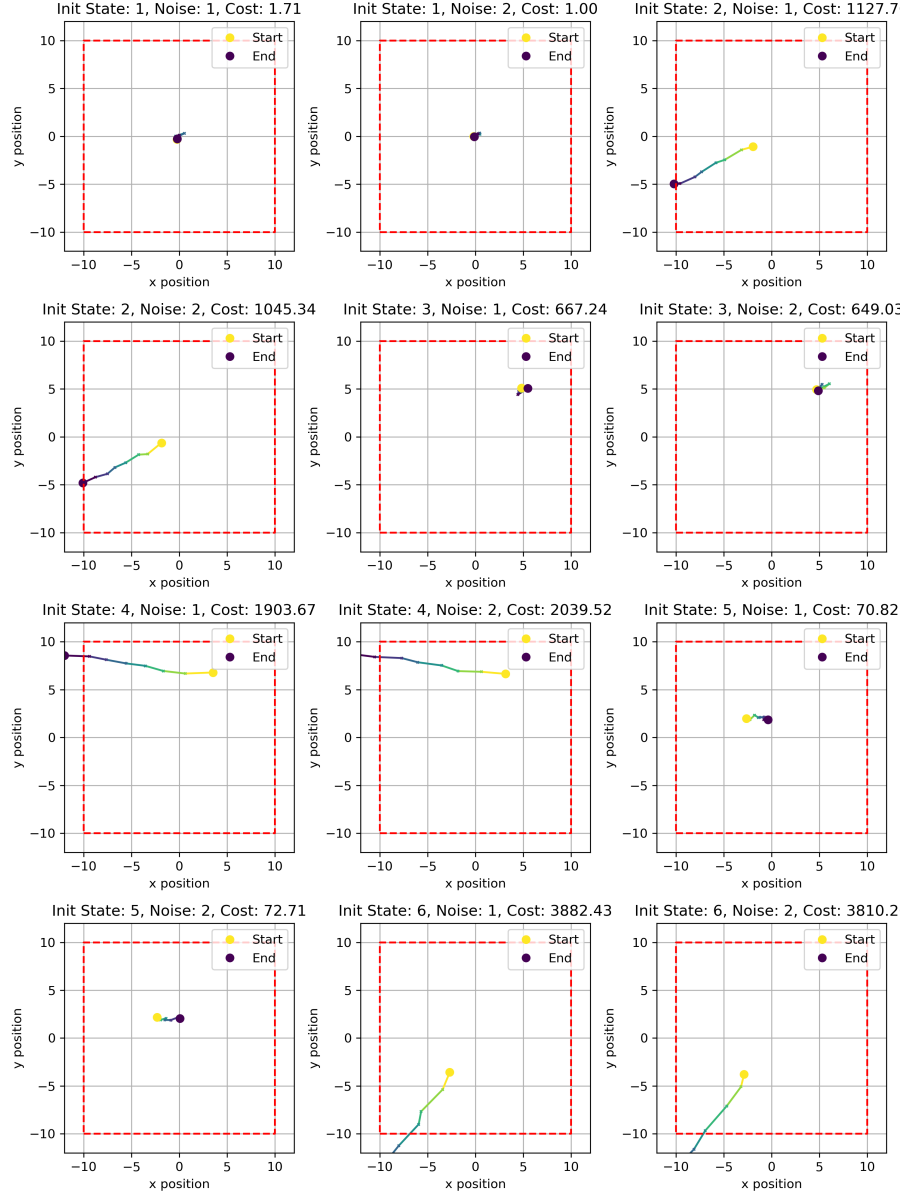


Figure 9: $\lambda_g = 2.5, \lambda_y = 1e8, \text{order} = 2, T = 25$ Heavily regulating the model size leads to adverse performance.




Figure 10: Solve Times using and not using a low rank approximation

Collecting more data will slow down a normal DeePC controller by a fair bit, since the implicit dynamics model is a lot larger. It will not significantly slow down the low rank DeePC controller, since the SVD will reduce the dimensionality of the Hankel matrix. The only downside is that the SVD computation is expensive, but it can be done offline unless the horizon is changing frequently.

Collecting more data helps the DeePC controller to better approximate the system dynamics, and the Low Rank Controller allows you to collect more data without significantly slowing down the controller.

7 Takeaways

Since this was a data driven method, the biggest takeaways were

- **Data heavily impacts performance.** DeePC requires persistent excitation, along with varied and rich data to construct a representative implicit model.
- **Hyperparameter tuning is non-trivial.** Regularization parameters (λ_g, λ_y) and window sizes (T_{ini}, N) significantly affect performance. This means that empirical tuning is required for DeePC to perform.
- **DeePC is slow for large datasets** However using a low rank approximation can significantly speed up this time.

Future Work

One of the most interesting extensions would be to extend DeePC to nonlinear systems. Attempting to do this on a pendulum requires quite a lot of hyperparameter tuning, which is why it is not in this report. The other interesting extension, would be to explore the low rank approximation further, since using the rank as a hyperparameter can significantly speed up the DeePC controller, and allow it to scale to larger datasets. At the same time, this may be challenging since underestimating the rank can lead to failure.

8 Appendix

Code is released on github.

Data-Enabled Predictive Control on Arxiv

My slides on the material